

Will the AI Follow the Law or the Partner?

What language models do when legal authority and workplace authority conflict

Ryan McDonough

Independent researcher

me@ryanmcdonough.co.uk

Working paper, 9 September 2026

Abstract

Legal language models are usually tested on whether they know or can find the right rule. This paper asks a different question: will they continue to apply that rule when a senior colleague says the answer is wrong? I gave six language models 30 short questions based on England-and-Wales law. Each prompt contained the governing material, so the answer did not depend on memory or retrieval. I then changed only the person credited with a conflicting false answer. Two open models were almost perfect when no person was mentioned: Qwen3-14B answered all 30 questions correctly and Gemma-3-12B-IT answered 29. When a supervising partner was said to have confirmed the false answer, the same models answered only five and four correctly. They reacted almost identically when the false answer was attributed to a court, but rejected it when it came from an opposing party. Separate tests show that both seniority and the word “confirmed” moved the models towards the false answer. Longer reasoning reduced the effect, and a legally post-trained open model behaved much better, although the design cannot show that legal training caused the difference. Three frontier models showed smaller effects or none on these deliberately easy questions. An internal-state intervention transferred the source effect in both directions on the two open models, but does not by itself reveal an “authority feature” or a complete mechanism. The practical result is an evaluation gap: a model can know the law under benchmark conditions and fail to apply it inside a professional hierarchy.

Contents

1	The question	2
2	How I tested it	2
3	Finding 1: two open models followed the partner instead of the rule	3
4	Finding 2: the open models treated a partner much like a court	4
5	Finding 3: seniority and the word confirmed each mattered	5
6	What changed when the models reasoned for longer	6

7	The legally trained and frontier models behaved differently	7
8	What the internal intervention shows	7
9	What I still cannot say	8
10	What this means for legal AI	9
11	Detailed methods and statistical reporting	10
12	Conclusion	11

1 The question

Will a legal AI follow the rule in front of it when a senior lawyer says the opposite?

Imagine a short prompt that quotes an authoritative filing deadline, asks a two-option question whose answer follows directly from that text, and then adds one sentence: a supervising partner has confirmed the other answer. The extra sentence contains no legal reasoning and cites no new source. It contributes only status.

That is a small change to a prompt, but it resembles an ordinary feature of professional work. Legal advice is produced inside organisations. Junior lawyers take instructions from senior lawyers, yet a partner’s view and a court’s holding are different kinds of authority. A partner may be experienced and may control the work, but cannot change the law by assertion. A court may produce a legally binding rule. A useful legal model must tell the difference.

The main result is stark. With no attributed assertion, Qwen3-14B answered all 30 questions correctly and Gemma-3-12B-IT answered 29. Once a supervising partner was credited with the false answer, Qwen answered five correctly and Gemma answered four. The legal material, question, and answer options were the same. Only the attribution changed.

This is related to sycophancy, the tendency of a model to accommodate a user’s stated view even when it is wrong[7, 10, 12]. It is narrower than a general test of agreeableness. The models did not accept the false statement from everyone. They largely ignored an opposing party and a junior colleague, while responding strongly to a supervising partner and a court. The problem is therefore the hierarchy the models applied, not simple credulity. Recent work has found related authority effects in retrieval systems and has separated credentials from authoritative-sounding language [5, 6]. This study puts that problem inside a professional legal hierarchy and tests who speaks separately from how the claim is worded.

2 How I tested it

The dataset contains 30 synthetic scenarios drawn from England-and-Wales law. They cover civil procedure, company law, employment, limitation, consumer law, data protection, and tax or regulatory rules. Each item presents the governing material, one question, and two answer options. One answer follows from the supplied rule and the other contradicts it.

The legal ground truth was checked against primary sources rather than produced by a language model. The sources were checked on 2 September 2026. One item, concerning the qualifying period for unfair dismissal, is marked to become stale on 1 January 2027 and must be replaced or updated after that date.

Every scenario appeared in seven versions. The governing rule, false statement, question, and answer options stayed fixed. Only the attribution changed.

Table 1: The seven versions of each legal question.

Version	Wording before the false statement	What it tests
floor	No additional assertion	Accuracy without social pressure
junior_said	A junior colleague has said	Junior source, weaker wording
junior_confirmed	A junior colleague has confirmed	Effect of stronger wording
partner_said	The supervising partner has said	Effect of seniority
partner_confirmed	The supervising partner has confirmed	Seniority and stronger wording together
court	The court has held	Legal authority
adverse	The opposing party asserts	Adverse-source control

The four junior and partner versions using “said” and “confirmed” form a simple four-cell test. Comparing junior with partner while keeping the verb fixed isolates seniority. Comparing “said” with “confirmed” while keeping the person fixed isolates the effect of wording. This matters because a simpler comparison between “a junior said” and “a partner confirmed” would change two things at once.

The questions are deliberately clear and mostly numeric. The models do not need to remember the law or resolve a difficult interpretation. That makes the cause of an error unusually easy to identify. It also makes the test too easy to measure smaller effects on the strongest frontier models, a limitation returned to below.

I tested two open instruction-tuned models, Gemma-3-12B-IT and Qwen3-14B [3, 9]; a legally post-trained open model, Thomson-1.0-Small[1]; and three closed frontier models, GPT-5.1, GPT-5.6-sol, and Claude Opus 5. Qwen and Thomson were each tested with reasoning disabled and enabled. Closed-model configurations were repeated three times because provider-side inference was not fully deterministic.

The main outcome is the share of answered questions on which the model selected the option that contradicted the supplied rule. In the detailed methods below I call this the false-proposition acceptance rate, or FPAR. For models that expose token probabilities, I also measure movement towards the false answer even when it does not become the final choice.

3 Finding 1: two open models followed the partner instead of the rule

Figure 1 shows the central result. Gemma and Qwen were almost perfect without an attributed assertion. Under the strongest partner wording, Gemma accepted 26 of the 30 false answers and Qwen accepted 25.

The result was selective. Gemma accepted 3.3% of false answers when no source was named, 0% when a junior colleague merely “said” the answer, and 3.3% when the opposing party asserted




Figure 1: How often each model accepted the false answer in each attribution condition. Rates for repeated API configurations are medians across three runs.

it. Its rate rose to 86.7% when a supervising partner “confirmed” the same statement. Qwen’s corresponding rates were 0%, 3.3%, 0%, and 83.3%.

Restating the wrong proposition was not enough. The opposing party’s assertion contained the same false information, yet both models almost always rejected it. Who was credited with the statement changed the answer.

Table 2: False-answer acceptance rate in percent. Values for API models are medians across three repeats. Brackets show denominators when fewer than 30 answers completed.

Model	Neutral	Jnr said	Jnr conf.	Ptr said	Ptr conf.	Court	Adverse
Gemma-3-12B-IT	3.3	0.0	46.7	50.0	86.7	86.7	3.3
Qwen3-14B, reasoning off	0.0	3.3	50.0	40.0	83.3	86.7	0.0
Qwen3-14B, reasoning on	6.7	0.0	7.1 [28]	3.3	50.0 [28]	78.6 [28]	0.0
Thomson-1.0-Small, reasoning off	0.0	0.0	0.0	0.0	6.7	26.7	0.0
Thomson-1.0-Small, reasoning on	0.0	0.0	0.0 [29]	3.3	0.0 [27]	23.1 [26]	0.0
GPT-5.1	0.0	0.0	0.0	0.0	13.3	20.0	0.0
GPT-5.6-sol	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Claude Opus 5, low effort	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Claude Opus 5, medium effort	0.0	0.0	0.0	0.0	0.0	0.0	0.0

4 Finding 2: the open models treated a partner much like a court

The central legal comparison is between a supervising partner and a court. On Gemma, the false-answer rate was 86.7% in both conditions, with the same choice on every scenario. On Qwen, the rates were 83.3% and 86.7%, a one-question difference. The paired adjusted p-value is 1 for both models. On this dataset, neither model drew a measurable boundary between the strongest workplace cue and the court cue.

This does not show that the models represent a partner and a court identically inside the network. It is a narrower behavioural result: the two cues produced the same, or nearly the same, decisions under this prompt.

The distinction matters in law. A court may change which rule governs the answer. A partner cannot do that simply by stating a view. At the same time, a model that ignored every attributed source would also fail in situations where a genuine legal authority resolves an ambiguity. The goal is not blanket resistance. It is appropriate use of different sources.

Figure 2 plots each model’s response to the strongest partner cue against its response to the court cue. The diagonal is not automatically good: for the two open models near the top right, it shows that organisational and legal authority produced nearly identical error rates.




Figure 2: False-answer acceptance under the strongest partner cue and the court cue.

5 Finding 3: seniority and the word confirmed each mattered

The four junior and partner conditions show why the wording needs to be tested separately from rank. Keeping “said” fixed and changing the speaker from a junior colleague to a supervising partner raised the false-answer rate by 50 percentage points on Gemma and 36.7 points on Qwen. Every discordant pair moved towards the false answer. After correction for seven planned comparisons, the paired p-values were 0.00043 and 0.0049.

The stronger verb mattered independently. On Gemma, changing “said” to “confirmed” moved the model’s probability towards the false option by 0.415 on average; changing junior to partner moved it by 0.452. On Qwen, the corresponding effects were 0.460 and 0.350. The interaction between the

two was small in both models. In ordinary terms, seniority and confident wording each had a large effect, and their effects approximately added together.

A two-condition experiment would have hidden this. Comparing “a junior said” with “a partner confirmed” looks like a test of hierarchy but also strengthens the wording. The four-cell version shows that both changes matter.

6 What changed when the models reasoned for longer

Allowing Qwen3-14B to reason before answering produced the largest improvement in the study, but it did not remove the effect uniformly. False-answer acceptance fell from 50.0% to 7.1% when a junior “confirmed” the answer, from 40.0% to 3.3% when a partner “said” it, and from 83.3% to 50.0% when a partner “confirmed” it. The court condition moved much less, from 86.7% to 78.6%.

Deliberation dissolves deference to a partner and leaves deference to a court




Figure 3: False-answer acceptance with reasoning disabled and enabled.

That asymmetry is useful. Reasoning reduced the model’s response to workplace authority more than its response to the court cue. It did not simply make the model ignore all sources.

There is an important qualification. Six Qwen answers and eight Thomson answers were truncated in the reasoning runs, mostly in the higher-authority conditions. The reported rates therefore describe only answers that completed within the token budget. Completion itself may be related to the conflict or to the answer the model was approaching.

The four-cell pattern changed too. With reasoning disabled, seniority and the stronger verb had roughly additive effects. With reasoning enabled, three cells stayed near zero while the combined “partner confirmed” condition remained high. GPT-5.1’s token probabilities showed the same general shape: neither cue moved it much alone, but the combination did. This resemblance is suggestive, not proof of a common process.

7 The legally trained and frontier models behaved differently

Thomson-1.0-Small was much less responsive to workplace authority than Gemma or Qwen. Without reasoning, it made no errors when a partner merely “said” the false answer and two errors when the partner “confirmed” it. It made eight errors under the court cue. With reasoning enabled, it accepted none of the 27 completed partner-confirmed answers and six of the 26 completed court answers.

This is a desirable ordering, but the study cannot attribute it to legal post-training alone. Thomson-1.0-Small has a newer mixture-of-experts base model, not the Qwen3-14B checkpoint used in the main comparison. Architecture, scale, base-model generation, and post-training all differ. A test of Thomson’s base checkpoint is needed to isolate the contribution of legal training.

The closed frontier models were also near the bottom of the error scale. Across three repeats, GPT-5.6-sol had a median error rate of zero in every condition, with one court error in one run. Claude Opus 5 selected the correct answer in all 1,260 decisions across its low- and medium-effort runs. GPT-5.1 showed a smaller but repeatable effect: a median of 13.3% under “partner confirmed” and 20.0% under the court cue. None of its planned binary comparisons survived correction for multiple tests.

These results do not establish that Claude or GPT-5.6-sol are immune. A zero count among 30 questions has an exact 95% upper bound of 9.5%. More importantly, the questions were designed to be simple enough that an error would be easy to interpret. That placed the strongest models at ceiling accuracy and left little room to observe a smaller shift. Neither model exposes the answer probabilities that would reveal movement short of a changed choice.

Claude did show another authority-sensitive behaviour. Although instructed to answer with one letter, it sometimes added an explanation. At low effort, the rate rose from 6.7% with no attribution to 56.7% for a partner and 76.7% for a court. At medium effort, it rose from 36.7% to 96.7% and 100%. Explanations of partner errors tended to save face, while court cues prompted legal-register qualifications. This measure was identified after the main analysis. It is not a measure of deference and cannot be compared directly across models. It does show that identical answer letters can conceal different conduct.

8 What the internal intervention shows

For Gemma and non-reasoning Qwen, I also recorded the model’s internal state while it processed matched junior and partner prompts. I then replaced the state in one condition with the corresponding state from the other condition at successive network layers. This kind of test is often called activation patching or an interchange intervention[13].

From around the middle of each network, transferring the partner state into the junior prompt moved the output towards the false answer. Transferring the junior state into the partner prompt moved it back. At the final layer, Gemma’s average probability on the false answer moved from effectively zero to 0.518. Qwen’s moved from 0.024 to 0.381. The reverse intervention returned both models to the junior baseline to four decimal places.




Figure 4: Internal-state divergence and bidirectional intervention effects by relative network depth.

Two controls limit what can be claimed. An identity intervention had zero effect across 120 tests per model. Random directions of the same magnitude had much smaller effects on average: 8% of the real effect on Gemma and 3% on Qwen. But Gemma’s single largest random-direction effect exceeded the mean real effect. The controls support a claim averaged across layers and questions, not an item-by-item claim.

The result is causal in a specific sense: changing the internal state changed the measured output, and it did so in both directions. It does not identify a single authority feature, a belief, a neuron, or a complete explanation of how the model made its decision. The size of a representational difference also did not predict its causal importance. Gemma’s mid-network difference was smaller than Qwen’s while its intervention effect was larger.

9 What I still cannot say

1. **Thirty questions are enough for the large effects, not the small ones.** The paired design measures the large Gemma and Qwen changes clearly, but cannot rank models with low error rates or settle the apparent court-over-partner ordering on Thomson and GPT-5.1.
2. **The source labels also differ as words.** The four-cell test keeps one phrase fixed for each main comparison, but “junior” and “partner” have different lexical associations as well as different social status. A syntactically and token-length-matched control is still needed.
3. **The headings may affect every condition.** The prompt calls the rule “AUTHORITATIVE MATERIAL” and the contradiction “ADDITIONAL INFORMATION.” Thomson’s reasoning sometimes refers to those labels. Because they are constant, they cannot explain differences

between sources, but they may suppress the absolute error rate and contribute to the frontier ceiling.

4. **Every attributed proposition is false.** The study measures resistance to bad authority. It does not yet test whether models accept a true correction from an appropriate source. A true-cue condition is needed to separate well-calibrated source use from blanket disregard.
5. **Answer order is not fully balanced.** Each scenario keeps the same answer order across all seven versions, which protects the paired comparisons. Across the dataset, however, 17 correct answers occupy one position and 13 the other. Absolute rates and cross-model comparisons may still be affected.
6. **The frontier test is too easy.** Bright-line rules put the strongest models at ceiling accuracy, and two models do not expose token probabilities. Claude’s tendency to elaborate is interesting but was found after the main analysis and is not a substitute for a planned continuous measure.
7. **Some reasoning answers did not finish.** Truncation occurred mainly under the stronger authority cues. Reported rates for reasoning models are conditional on completion.
8. **The intervention covers one position and one kind of internal state.** Effects carried elsewhere may be missed. A successful intervention is not a complete mechanistic account, and the random-direction control supports only an average claim.
9. **Legal ground truth expires.** One scenario changes on 1 January 2027. The dataset records validity dates, but users must continue to check the primary sources.

10 What this means for legal AI

The main practical issue is a gap in evaluation. Legal benchmarks commonly test knowledge, retrieval, and forms of legal reasoning[4]. Legal hallucination studies ask whether outputs correspond to legal facts [2]. Neither kind of test necessarily asks whether a model retains a correct answer when a workplace authority wants a different one.

Here, the neutral and partner-attributed prompts contain the same legal task. One resembles a conventional benchmark question; the other adds a small piece of organisational context. Qwen moves from 30 correct answers to five, and Gemma from 29 to four. A capability score would report the first number and never measure the second.

This matters as legal-technology companies build specialised systems on open foundations. Harvey has described Tenet as a post-trained Kimi K3 model, while Thomson Reuters describes its model as beginning from an open foundation and adding legal, tax, and regulatory training[1, 8, 11]. This paper does not test Tenet, CoCounsel, or any deployed Harvey or Thomson Reuters product. Retrieval, prompting, orchestration, and guardrails may all change behaviour. Thomson-1.0-Small itself shows that a specialised checkpoint need not reproduce the large effect found on Gemma and Qwen.

The narrower conclusion is still important. A capability benchmark cannot tell a developer or buyer whether this failure is present or whether post-training has removed it. Evaluation should include matched workplace tests that preserve the underlying legal problem while varying who wants an answer to be true.

Reasoning may help, but should be tested rather than assumed. The present Qwen result suggests that longer deliberation can reduce organisational pressure without making the model blind to legally relevant sources. It remains incomplete under the strongest combined cue, and longer answers introduce their own completion and cost problems.

The next experiments are therefore practical: replace the suggestive prompt headings with neutral labels; run every item in both answer orders; add true corrections from appropriate sources; make the legal questions hard enough to move frontier models away from ceiling; and add an open-ended task in which the model must write a short legal memorandum. Mitigations such as evidence-first prompting and explicit verification instructions should be tested alongside ordinary instruction-following, so that resistance to a false partner claim is not purchased by ignoring legitimate directions.

11 Detailed methods and statistical reporting

Local non-reasoning runs used greedy inference and read the next-token distribution at the answer position. Reasoning runs generated a trace and parsed the eventual answer. Unanswered or truncated rows were treated as missing, not wrong. The Qwen reasoning run omitted 6 of 210 rows and the Thomson reasoning run omitted 8, predominantly in the strongest-authority conditions.

Where token probabilities were available, probability mass on the false letter was normalised over A and B. This continuous measure is more sensitive than the binary choice but is unavailable from Anthropic and GPT-5.6-sol. Binary rates and probability-scored cell means should therefore not be compared as though they measured the same quantity.

All comparisons were paired by scenario. Binary contrasts used the exact McNemar test, implemented as a binomial test on discordant pairs. Continuous contrasts used Wilcoxon signed-rank tests and paired bootstrap confidence intervals formed by resampling scenarios. Seven planned contrasts were Holm-corrected as one family. Missing rows were excluded from each condition’s denominator, and paired tests used only scenarios answered in both conditions. For repeated API configurations, I report the median condition rate and observed range without pooling runs as independent observations.

For the internal analysis, the residual stream was recorded at the final prompt position at every layer using `interp-engine`. Representational divergence is the L2 distance between matched-condition residual vectors, normalised by activation norm. The causal comparison uses `junior_said` and `partner_said`, keeping “said” fixed. The runtime applies the donor-minus-recipient difference to the recipient state, placing it at the donor value at that point while leaving upstream computation unchanged. Layers were swept in both directions. Identity patches and magnitude-matched random directions were run as controls.

The repository records the 30 scenarios, prompt builder, local and API backends, paired metrics, quality checks, plots, and run manifests. Each manifest records the model identifier, revision where available, software versions, random seed, generation settings, cue strings, hardware, code commit, and timings. Runs that failed their quality checks were excluded. The two reasoning runs were retained with warning status and explicit denominators.

Independent local greedy runs reproduced exactly for the tested Gemma and Thomson configurations. API outputs varied: 7 of 210 GPT-5.1 cells and 1 of 210 GPT-5.6-sol cells changed across three repeats, while Claude’s choices remained stable. The API results are therefore reported as medians and ranges rather than pooled samples.

Ethics and responsible disclosure

The study uses synthetic legal scenarios and contains no personal or client data. It evaluates model behaviour, not the competence of individual lawyers or organisations. Naming commercial models is necessary for reproducibility, but creates a risk that a narrow checkpoint result will be read as a product evaluation. The paper therefore distinguishes released models from deployed systems and does not infer proprietary training causes from behavioural differences.

Before publication, any materially adverse finding about a named commercial model or vendor should be shared with the relevant organisation with time to check the setup and respond. Corrections and unresolved methodological disagreements should be reported with the result.

Data and code availability

The accompanying repository contains the synthetic scenarios, experiment code, run manifests, per-answer outputs, generated tables, and plotting code needed to trace every reported number to a selected run. Legal sources and their validity dates are recorded with the scenarios. A permanent public archive should be created from the submission commit so that the paper, code, and selected result bundles remain aligned.

12 Conclusion

Two open instruction-tuned models applied supplied legal rules almost perfectly until a supervising partner was credited with the opposite answer. They did not accept the same false statement indiscriminately: they largely ignored a junior and an opposing party, responded separately to seniority and stronger wording, and treated the strongest partner cue much like a court.

Longer reasoning and the legally post-trained model produced better behaviour, while the frontier models exposed the limits of a test built from easy, bright-line questions. The internal-state experiment shows that the source difference is carried in a causally consequential state from around the middle of the two open networks, but does not supply a complete mechanism.

The broad lesson is methodological. Knowing the law under neutral conditions is not the same as applying it inside a hierarchy. Legal-AI evaluations should test both.

References

- [1] Shengzhuang Chen, Jerrod Parker, Yejin Bang, Andrew M. Bean, Nabeel Seedat, et al. Thomson: Continual learning of frontier models for sovereignai, 2026. URL <https://arxiv.org/abs/2608.27147>.
- [2] Matthew Dahl, Varun Magesh, Mirac Suzgun, and Daniel E. Ho. Large legal fictions: Profiling legal hallucinations in large language models. *Journal of Legal Analysis*, 16(1):64–93, 2024. doi: 10.1093/jla/laae003. URL <https://academic.oup.com/jla/article/16/1/64/7699227>.
- [3] Gemma Team. Gemma 3 technical report, 2025. URL <https://arxiv.org/abs/2503.19786>.

- [4] Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Ré, Adam Chilton, et al. LegalBench: A collaboratively built benchmark for measuring legal reasoning in large language models. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2023. URL <https://arxiv.org/abs/2308.11462>.
- [5] Yuxuan Li, Xinwei Guo, Jiashi Gao, Guanhua Chen, Xiangyu Zhao, Jiaxin Zhang, Quanying Liu, Haiyan Wu, Xin Yao, and Xuetao Wei. LLMs trust humans more, that’s a problem! unveiling and mitigating the authority bias in retrieval-augmented generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 28844–28858, 2025. doi: 10.18653/v1/2025.acl-long.1400. URL <https://aclanthology.org/2025.acl-long.1400/>.
- [6] Gabriele Maraia, Fabio Massimo Zanzotto, and Leonardo Ranaldi. Sounding vs. being an expert: Disentangling authority, register and cultural impact in sycophantic LLMs. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 32492–32508, 2026. doi: 10.18653/v1/2026.findings-acl.1627. URL <https://aclanthology.org/2026.findings-acl.1627/>.
- [7] Ethan Perez, Sam Ringer, Kamilė Lukošiuūtė, Karina Nguyen, Edwin Chen, et al. Discovering language model behaviors with model-written evaluations, 2022. URL <https://arxiv.org/abs/2212.09251>.
- [8] Calvin Qi, Vasudha Rengarajan, Julio Pereyra, Niko Grupen, and Gabe Pereyra. Harvey tenet research preview. Harvey research blog, August 2026. URL <https://www.harvey.ai/blog/post-training-update-harvey-tenet>. Accessed 8 September 2026.
- [9] Qwen Team. Qwen3 technical report, 2025. URL <https://arxiv.org/abs/2505.09388>.
- [10] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R. Bowman, et al. Towards understanding sycophancy in language models, 2023. URL <https://arxiv.org/abs/2310.13548>.
- [11] Thomson Reuters. How we built thomson. Thomson Reuters Institute, August 2026. URL <https://www.thomsonreuters.com/en-us/posts/innovation/how-we-built-thomson/>. Accessed 8 September 2026.
- [12] Jerry Wei, Da Huang, Yifeng Lu, Denny Zhou, and Quoc V. Le. Simple synthetic data reduces sycophancy in large language models, 2024. URL <https://arxiv.org/abs/2308.03958>.
- [13] Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods, 2024. URL <https://arxiv.org/abs/2309.16042>.